

Starting the Cluster

Task 1

Task1: Retail Store Insights

Scenario: You're a data analyst at a large retail store. The store sells a variety of products, including books and fruits. The management wants insights into sales patterns, customer preferences, product popularity, and potential promotions.

Objective: Analyze the provided dataset to extract meaningful insights and present them to the management.

Instructions:

1. Initialization:

- Set up your environment by initializing a Spark session. Name this session "RetailStoreInsights".

In [1]:

```
1 from pyspark.sql import SparkSession
2
3 spark = SparkSession.builder \
4     .remote("sc://192.168.1.7:15002") \
5     .appName("RetailStoreInsights") \
6     .getOrCreate()
7
8 # limit() shows a nice HTML table in Jupyter, while show() prints plain text
9 spark.conf.set('spark.sql.repl.eagerEval.enabled', True)
10
11 spark
```

Python

Out[1]:

```
1 <pyspark.sql.connect.session.SparkSession at 0x29256293a10>
```

out

In [2]:

```
1 from pyspark.sql import Row
2
3 data = [
4     ('Ulysses', 'Book', 23.17, 16),
5     ('Apple', 'Fruit', 2.34, 8),
6     ('Pineapple', 'Fruit', 2.57, 1),
7     ('Apple', 'Fruit', 2.43, 6),
8     ('To Kill a Mockingbird', 'Book', 24.14, 19),
9     ('To Kill a Mockingbird', 'Book', 11.18, 11),
10    ('Watermelon', 'Fruit', 3.35, 15),
11    ('Pride and Prejudice', 'Book', 24.99, 3),
12    ('To Kill a Mockingbird', 'Book', 21.82, 17),
13    ('Moby Dick', 'Book', 14.83, 20),
14    ('Pride and Prejudice', 'Book', 5.03, 16),
15    ('Jane Eyre', 'Book', 20.40, 8),
```

Python

```

16 ('Moby Dick', 'Book', 5.55, 20),
17 ('Don Quixote', 'Book', 19.75, 17),
18 ('Watermelon', 'Fruit', 2.31, 9),
19 ('Hamlet', 'Book', 18.20, 12),
20 ('Mango', 'Fruit', 4.10, 7),
21 ('1984', 'Book', 16.75, 14),
22 ('Strawberry', 'Fruit', 1.90, 25),
23 ('War and Peace', 'Book', 22.50, 9),
24 ('Orange', 'Fruit', 3.05, 13),
25 ('The Great Gatsby', 'Book', 12.30, 10),
26 ('Peach', 'Fruit', 2.80, 11),
27 ('Grapes', 'Fruit', 2.60, 18),
28 ('Pride and Prejudice', 'Book', 9.50, 5)
29 ]
30
31 df = spark.createDataFrame([
32     Row(product_name=row[0], category=row[1], price=row[2], quantity=row[3])
33     for row in data
34 ], schema = 'product_name STRING, category STRING, price FLOAT, quantity
SHORT')
35 df.createOrReplaceTempView("retail_sales") # give it a name for sql
36 df.limit(10)

```

Out[2]:

product_name	category	price	quantity
Ulysses	Book	23.17	16
Apple	Fruit	2.34	8
Pineapple	Fruit	2.57	1
Apple	Fruit	2.43	6
To Kill a Mocking...	Book	24.14	19
To Kill a Mocking...	Book	11.18	11
Watermelon	Fruit	3.35	15
Pride and Prejudice	Book	24.99	3
To Kill a Mocking...	Book	21.82	17
Moby Dick	Book	14.83	20

out

In [3]:

```
1 df.printSchema()
```

Python

```

1 root
2 |-- product_name: string (nullable = true)
3 |-- category: string (nullable = true)

```

out

```

4 |-- price: float (nullable = true)
5 |-- quantity: short (nullable = true)
6


```

In [4]:

```

1 spark.sql("""
2   select * from retail_sales
3   where price > 2
4   order by price
5   """)

```

 Python

Out[4]:

```

1  +-----+-----+-----+-----+
2  |      product_name|category|price|quantity|
3  +-----+-----+-----+-----+
4  |      Watermelon|   Fruit| 2.31|       9|
5  |           Apple|   Fruit| 2.34|       8|
6  |           Apple|   Fruit| 2.43|       6|
7  |      Pineapple|   Fruit| 2.57|       1|
8  |           Grapes|   Fruit|  2.6|      18|
9  |           Peach|   Fruit|  2.8|      11|
10 |           Orange|   Fruit| 3.05|      13|
11 |      Watermelon|   Fruit| 3.35|      15|
12 |           Mango|   Fruit|  4.1|       7|
13 |Pride and Prejudice|   Book| 5.03|      16|
14 |           Moby Dick|   Book| 5.55|      20|
15 |Pride and Prejudice|   Book|  9.5|       5|
16 |To Kill a Mocking...|   Book|11.18|      11|
17 |   The Great Gatsby|   Book| 12.3|      10|
18 |           Moby Dick|   Book|14.83|      20|
19 |           1984|   Book|16.75|      14|
20 |           Hamlet|   Book| 18.2|      12|
21 |      Don Quixote|   Book|19.75|      17|
22 |           Jane Eyre|   Book| 20.4|       8|
23 |To Kill a Mocking...|   Book|21.82|      17|
24 +-----+-----+-----+-----+
25 only showing top 20 rows

```

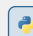
out

In [5]:

```

1 # https://spark.apache.org/docs/latest/sql-pipe-syntax.html
2 spark.sql("""
3   from retail_sales
4   |> aggregate count(*) as category_count
5   group by category
6   """)

```

 Python

Out[5]:

```

1 +-----+-----+
2 |category|category_count|
3 +-----+-----+
4 |   Book|           15|
5 |   Fruit|           10|
6 +-----+-----+

```

out

In [6]:

```

1 spark.sql("""
2   from retail_sales
3   |> aggregate avg(price) as avg_price
4     group by product_name
5   |> set avg_price = round(avg_price, 2)
6   """)

```

Python

Out[6]:

```

1 +-----+-----+
2 |      product_name|avg_price|
3 +-----+-----+
4 |      Pineapple|      2.57|
5 |To Kill a Mocking...|     19.05|
6 |      Ulysses|     23.17|
7 |      Apple|      2.38|
8 |      Jane Eyre|      20.4|
9 |      Moby Dick|     10.19|
10 |      Watermelon|      2.83|
11 |Pride and Prejudice|     13.17|
12 |      1984|     16.75|
13 |      Mango|       4.1|
14 |      Don Quixote|     19.75|
15 |      Hamlet|      18.2|
16 |      Orange|      3.05|
17 |      Peach|       2.8|
18 |The Great Gatsby|     12.3|
19 |      Grapes|       2.6|
20 |      Strawberry|       1.9|
21 |War and Peace|     22.5|
22 +-----+-----+

```

out

In [7]:

```

1 spark.sql("""
2   from retail_sales
3   |> extend price - (price * 0.1) as discounted_price
4   |> set discounted_price = round(discounted_price, 2)
5   |> select product_name, discounted_price, price as original_price
6   """)

```

Python

Out[7]:

1	+-----+-----+-----+			out
2		product_name discounted_price original_price		
3	+-----+-----+-----+			
4		Ulysses 20.85 23.17		
5		Apple 2.11 2.34		
6		Pineapple 2.31 2.57		
7		Apple 2.19 2.43		
8		To Kill a Mocking... 21.73 24.14		
9		To Kill a Mocking... 10.06 11.18		
10		Watermelon 3.01 3.35		
11		Pride and Prejudice 22.49 24.99		
12		To Kill a Mocking... 19.64 21.82		
13		Moby Dick 13.35 14.83		
14		Pride and Prejudice 4.53 5.03		
15		Jane Eyre 18.36 20.4		
16		Moby Dick 5.0 5.55		
17		Don Quixote 17.78 19.75		
18		Watermelon 2.08 2.31		
19		Hamlet 16.38 18.2		
20		Mango 3.69 4.1		
21		1984 15.08 16.75		
22		Strawberry 1.71 1.9		
23		War and Peace 20.25 22.5		
24	+-----+-----+-----+			
25	only showing top 20 rows			

In [8]:

```
1 spark.sql("""
2   from retail_sales
3   |> aggregate sum(quantity) as n_sold_total
4   """)
```

Python

Out[8]:

1	+-----+			out
2		n_sold_total		
3	+-----+			
4		310		
5	+-----+			

In [9]:

```
1 spark.sql("""
2   from retail_sales
3   |> aggregate sum(quantity) as n_sold
4   group by category
5   """)
```

Python

Out[9]:

```

1 +-----+-----+
2 |category|n_sold|
3 +-----+-----+
4 |   Book|   197|
5 |   Fruit|   113|
6 +-----+-----+

```

out

In [10]:

```

1 spark.sql("""
2   from retail_sales
3   |> aggregate sum(price * quantity) as revenue
4   group by category
5   """)

```

Python

Out[10]:

```

1 +-----+-----+
2 |category|          revenue|
3 +-----+-----+
4 |   Book|3211.2000007629395|
5 |   Fruit| 300.3599935770035|
6 +-----+-----+

```

out

In [11]:

```

1 spark.sql("""
2   from retail_sales
3   |> aggregate sum(quantity) as n_sold
4   group by category, product_name
5   |> order by n_sold desc
6   """)

```

Python

Out[11]:

```

1 +-----+-----+-----+
2 |category| product_name|n_sold|
3 +-----+-----+-----+
4 |   Book|To Kill a Mocking...|   47|
5 |   Book|      Moby Dick|   40|
6 |   Fruit|      Strawberry|   25|
7 |   Fruit|      Watermelon|   24|
8 |   Book|Pride and Prejudice|   24|
9 |   Fruit|          Grapes|   18|
10 |   Book|      Don Quixote|   17|
11 |   Book|          Ulysses|   16|
12 |   Fruit|          Apple|   14|
13 |   Book|          1984|   14|
14 |   Fruit|          Orange|   13|
15 |   Book|          Hamlet|   12|
16 |   Fruit|          Peach|   11|

```


out

17		Book	The Great Gatsby	10
18		Book	War and Peace	9
19		Book	Jane Eyre	8
20		Fruit	Mango	7
21		Fruit	Pineapple	1
22	+-----+-----+-----+			

teste

In []:


```
1 spark.stop()
```

 Python

Task 2

In [1]:

```
1 from pyspark.sql import SparkSession
2
3 spark = SparkSession.builder \
4     .remote("sc://192.168.1.7:15002") \
5     .appName("UDFTransformation") \
6     .config("spark.sql.ansi.enabled", "false") \
7     .config("spark.sql.execution.pythonUDF.arrow.enabled", "true") \
8     .getOrCreate()
9
10 # limit() shows a nice HTML table in Jupyter, while show() prints plain text
11 spark.conf.set('spark.sql.repl.eagerEval.enabled', True)
12
13 spark
```

 Python


Out[1]:

```
1 <pyspark.sql.connect.session.SparkSession at 0x23f824ef550>
```

out

In [2]:

```
1 from pyspark.sql.functions import udf
2
3 @udf(returnType='int')
4 def mult_by_3(s: int) -> int:
5     return s * 3
6
7 df = spark.createDataFrame([(4, ), (5, ), (6, )], ['value'])
8 df
```

 Python


Out[2]:

```
1 +-----+
2 |value|
3 +-----+
4 |    4|
5 |    5|
6 |    6|
7 +-----+
```

out

In [3]:

```
1 dff = df.withColumn('value_x3', mult_by_3(df.value))
2 dff
```

 Python

Out[3]:


```
1 +-----+-----+
2 |value|value_x3|
3 +-----+-----+
4 |    4|      12|
```

out

5		5		15	
6		6		18	
7	+-----+-----+				

In [4]:

```
1 import pandas as pd
2 import pyspark.pandas as ps
3 from pyspark.sql.functions import pandas_udf
4
5 @pandas_udf("int")
6 def sub_2(s: pd.Series) -> pd.Series:
7     return s - 2
8
9 dffs = dff.withColumn('value_minus_2', sub_2(dff.value))
10 dffs
```

 Python

```
c:\Users\plancha\spark-lab1\.venv\lib\site-
packages\pyspark\pandas\__init__.py:43: UserWarning:
1 'PYARROW_IGNORE_TIMEZONE' environment variable was not set. It is required
  to set this environment variable to '1' in both driver and executor sides
  if you use pyarrow>=2.0.0. pandas-on-Spark will set it for you but it does
  not work if there is a Spark context already launched.
2 warnings.warn(
```

out


Out[4]:

1	+-----+-----+-----+				
2	value value_x3 value_minus_2				
3	+-----+-----+-----+				
4		4		12	
5		5		15	
6		6		18	
7	+-----+-----+-----+				

out

In []:

```
1 spark.stop()
```

 Python