

Task 1

In [1]: `from pyspark.sql import SparkSession`

```
spark = SparkSession.builder \
    .remote("sc://192.168.1.7:15002") \
    .appName("RetailStoreInsights") \
    .getOrCreate()

# limit() shows a nice HTML table in Jupyter, while show() prints plain text
spark.conf.set('spark.sql.repl.eagerEval.enabled', True)

spark
```

Out[1]: `<pyspark.sql.connect.session.SparkSession at 0x29256293a10>`

In [2]: `from pyspark.sql import Row`

```
data = [
    ('Ulysses', 'Book', 23.17, 16),
    ('Apple', 'Fruit', 2.34, 8),
    ('Pineapple', 'Fruit', 2.57, 1),
    ('Apple', 'Fruit', 2.43, 6),
    ('To Kill a Mockingbird', 'Book', 24.14, 19),
    ('To Kill a Mockingbird', 'Book', 11.18, 11),
    ('Watermelon', 'Fruit', 3.35, 15),
    ('Pride and Prejudice', 'Book', 24.99, 3),
    ('To Kill a Mockingbird', 'Book', 21.82, 17),
    ('Moby Dick', 'Book', 14.83, 20),
    ('Pride and Prejudice', 'Book', 5.03, 16),
    ('Jane Eyre', 'Book', 20.40, 8),
    ('Moby Dick', 'Book', 5.55, 20),
    ('Don Quixote', 'Book', 19.75, 17),
    ('Watermelon', 'Fruit', 2.31, 9),
    ('Hamlet', 'Book', 18.20, 12),
    ('Mango', 'Fruit', 4.10, 7),
    ('1984', 'Book', 16.75, 14),
    ('Strawberry', 'Fruit', 1.90, 25),
    ('War and Peace', 'Book', 22.50, 9),
    ('Orange', 'Fruit', 3.05, 13),
    ('The Great Gatsby', 'Book', 12.30, 10),
    ('Peach', 'Fruit', 2.80, 11),
    ('Grapes', 'Fruit', 2.60, 18),
    ('Pride and Prejudice', 'Book', 9.50, 5)
]

df = spark.createDataFrame([
    Row(product_name=row[0], category=row[1], price=row[2], quantity=row[3])
    for row in data
], schema = 'product_name STRING, category STRING, price FLOAT, quantity SHORT')
df.createOrReplaceTempView("retail_sales") # give it a name for sql
df.limit(10)
```

```
Out[2]: +-----+-----+-----+-----+
|      product_name|category|price|quantity|
+-----+-----+-----+-----+
|          Ulysses|    Book|23.17|      16|
|           Apple|   Fruit| 2.34|       8|
|        Pineapple|   Fruit| 2.57|       1|
|           Apple|   Fruit| 2.43|       6|
|To Kill a Mocking...|    Book|24.14|      19|
|To Kill a Mocking...|    Book|11.18|      11|
|        Watermelon|   Fruit| 3.35|      15|
|Pride and Prejudice|    Book|24.99|       3|
|To Kill a Mocking...|    Book|21.82|      17|
|         Moby Dick|    Book|14.83|      20|
+-----+-----+-----+-----+
```

```
In [3]: df.printSchema()

root
 |-- product_name: string (nullable = true)
 |-- category: string (nullable = true)
 |-- price: float (nullable = true)
 |-- quantity: short (nullable = true)
```

```
In [4]: spark.sql("""
        select * from retail_sales
        where price > 2
        order by price
        """)
```

```
Out[4]: +-----+-----+-----+-----+
|      product_name|category|price|quantity|
+-----+-----+-----+-----+
|        Watermelon|   Fruit| 2.31|       9|
|           Apple|   Fruit| 2.34|       8|
|           Apple|   Fruit| 2.43|       6|
|        Pineapple|   Fruit| 2.57|       1|
|          Grapes|   Fruit| 2.6|      18|
|          Peach|   Fruit| 2.8|      11|
|          Orange|   Fruit| 3.05|      13|
|        Watermelon|   Fruit| 3.35|      15|
|          Mango|   Fruit| 4.1|       7|
|Pride and Prejudice|    Book| 5.03|      16|
|          Moby Dick|    Book| 5.55|      20|
|Pride and Prejudice|    Book| 9.5|       5|
|To Kill a Mocking...|    Book|11.18|      11|
|    The Great Gatsby|    Book|12.3|      10|
|          Moby Dick|    Book|14.83|      20|
|          1984|    Book|16.75|      14|
|          Hamlet|    Book|18.2|      12|
|        Don Quixote|    Book|19.75|      17|
|          Jane Eyre|    Book|20.4|       8|
|To Kill a Mocking...|    Book|21.82|      17|
```

```
+-----+-----+-----+-----+
only showing top 20 rows
```

```
In [5]: # https://spark.apache.org/docs/latest/sql-pipe-syntax.html
spark.sql("""
    from retail_sales
    |> aggregate count(*) as category_count
    group by category
""")
```

```
Out[5]: +-----+-----+
|category|category_count|
+-----+-----+
|    Book|             15|
|   Fruit|             10|
+-----+-----+
```

```
In [6]: spark.sql("""
    from retail_sales
    |> aggregate avg(price) as avg_price
    group by product_name
    |> set avg_price = round(avg_price, 2)
""")
```

```
Out[6]: +-----+-----+
|      product_name|avg_price|
+-----+-----+
|      Pineapple|    2.57|
|To Kill a Mocking...|   19.05|
|      Ulysses|   23.17|
|      Apple|    2.38|
|      Jane Eyre|   20.4|
|      Moby Dick|   10.19|
|    Watermelon|    2.83|
|Pride and Prejudice|   13.17|
|      1984|   16.75|
|      Mango|    4.1|
|    Don Quixote|   19.75|
|      Hamlet|   18.2|
|      Orange|    3.05|
|      Peach|    2.8|
|    The Great Gatsby|   12.3|
|      Grapes|    2.6|
|    Strawberry|    1.9|
|    War and Peace|   22.5|
+-----+-----+
```

```
In [7]: spark.sql("""
    from retail_sales
    |> extend price - (price * 0.1) as discounted_price
    |> set discounted_price = round(discounted_price, 2)
    |> select product_name, discounted_price, price as original_price
""")
```

```
Out[7]: +-----+-----+-----+
|      product_name|discounted_price|original_price|
+-----+-----+-----+
|      Ulysses|      20.85|      23.17|
|      Apple|      2.11|      2.34|
|    Pineapple|      2.31|      2.57|
|      Apple|      2.19|      2.43|
|To Kill a Mocking...|      21.73|      24.14|
|To Kill a Mocking...|      10.06|      11.18|
|    Watermelon|      3.01|      3.35|
|Pride and Prejudice|      22.49|      24.99|
|To Kill a Mocking...|      19.64|      21.82|
|    Moby Dick|      13.35|      14.83|
|Pride and Prejudice|      4.53|      5.03|
|    Jane Eyre|      18.36|      20.4|
|    Moby Dick|      5.0|      5.55|
|    Don Quixote|      17.78|      19.75|
|    Watermelon|      2.08|      2.31|
|    Hamlet|      16.38|      18.2|
|    Mango|      3.69|      4.1|
|    1984|      15.08|      16.75|
|    Strawberry|      1.71|      1.9|
|    War and Peace|      20.25|      22.5|
+-----+-----+-----+
only showing top 20 rows
```

```
In [8]: spark.sql("""
        from retail_sales
        |> aggregate sum(quantity) as n_sold_total
        """)
```

```
Out[8]: +-----+
|n_sold_total|
+-----+
|      310|
+-----+
```

```
In [9]: spark.sql("""
        from retail_sales
        |> aggregate sum(quantity) as n_sold
        group by category
        """)
```

```
Out[9]: +-----+-----+
|category|n_sold|
+-----+-----+
|    Book|    197|
|    Fruit|    113|
+-----+-----+
```

```
In [10]: spark.sql("""
        from retail_sales
        |> aggregate sum(price * quantity) as revenue
```

```
group by category
""")
```

```
Out[10]: +-----+-----+
|category|      revenue|
+-----+-----+
|   Book|3211.2000007629395|
|   Fruit| 300.3599935770035|
+-----+-----+
```

```
In [11]: spark.sql("""
    from retail_sales
    |> aggregate sum(quantity) as n_sold
    group by category, product_name
    |> order by n_sold desc
    """)
```

```
Out[11]: +-----+-----+-----+
|category|      product_name|n_sold|
+-----+-----+-----+
|   Book|To Kill a Mocking...|   47|
|   Book|      Moby Dick|   40|
|   Fruit|      Strawberry|   25|
|   Fruit|      Watermelon|   24|
|   Book|Pride and Prejudice|   24|
|   Fruit|      Grapes|   18|
|   Book|      Don Quixote|   17|
|   Book|      Ulysses|   16|
|   Fruit|      Apple|   14|
|   Book|      1984|   14|
|   Fruit|      Orange|   13|
|   Book|      Hamlet|   12|
|   Fruit|      Peach|   11|
|   Book|The Great Gatsby|   10|
|   Book|War and Peace|    9|
|   Book|Jane Eyre|    8|
|   Fruit|      Mango|    7|
|   Fruit|Pineapple|    1|
+-----+-----+-----+
```

teste

```
In [ ]: spark.stop()
```

Task 2

```
In [1]: from pyspark.sql import SparkSession
```

```

spark = SparkSession.builder \
    .remote("sc://192.168.1.7:15002") \
    .appName("UDFTransformation") \
    .config("spark.sql.ansi.enabled", "false") \
    .config("spark.sql.execution.pythonUDF.arrow.enabled", "true") \
    .getOrCreate()

# limit() shows a nice HTML table in Jupyter, while show() prints plain text
spark.conf.set('spark.sql.repl.eagerEval.enabled', True)

spark

```

Out[1]: <pyspark.sql.connect.session.SparkSession at 0x23f824ef550>

```

In [2]: from pyspark.sql.functions import udf

@udf(returnType='int')
def mult_by_3(s: int) -> int:
    return s * 3

df = spark.createDataFrame([(4, ), (5, ), (6, )], ['value'])
df

```

Out[2]:

value
4
5
6

```

In [3]: dff = df.withColumn('value_x3', mult_by_3(df.value))
dff

```

Out[3]:

value	value_x3
4	12
5	15
6	18

```

In [4]: import pandas as pd
import pyspark.pandas as ps
from pyspark.sql.functions import pandas_udf

@pandas_udf("int")
def sub_2(s: pd.Series) -> pd.Series:
    return s - 2

dffs = dff.withColumn('value_minus_2', sub_2(dff.value))
dffs

```

```
c:\Users\plancha\spark-lab1\.venv\lib\site-packages\pyspark\pandas\__init__.py:43:
UserWarning: 'PYARROW_IGNORE_TIMEZONE' environment variable was not set. It is
required to set this environment variable to '1' in both driver and executor sides
if you use pyarrow>=2.0.0. pandas-on-Spark will set it for you but it does not work
if there is a Spark context already launched.
  warnings.warn(
```

```
Out[4]: +-----+-----+-----+
|value|value_x3|value_minus_2|
+-----+-----+-----+
|   4|      12|           2|
|   5|      15|           3|
|   6|      18|           4|
+-----+-----+-----+
```

```
In [ ]: spark.stop()
```